

AWS training: BoTorch TorchX - Neural architecture search

EC2 setup (free tier)

Credentials

Create job queue

Set user permissions

Python code

Installation

mnist.py

Jupyter notebook

AWS credentials

TorchXRunner

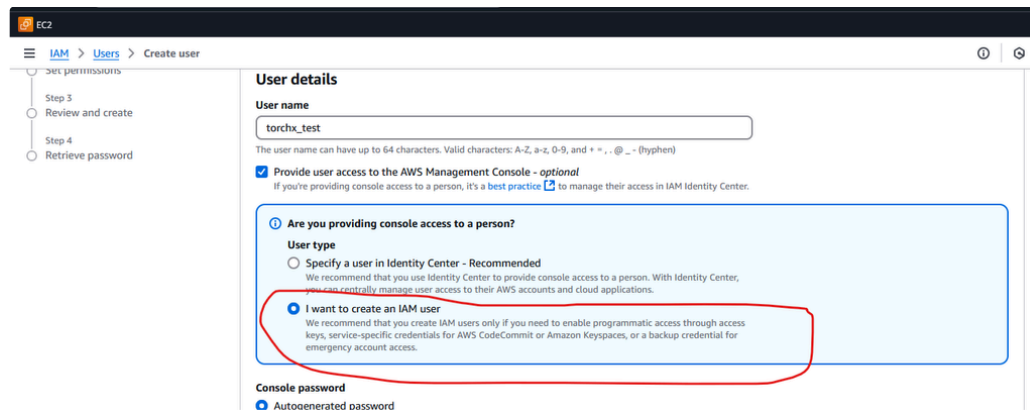
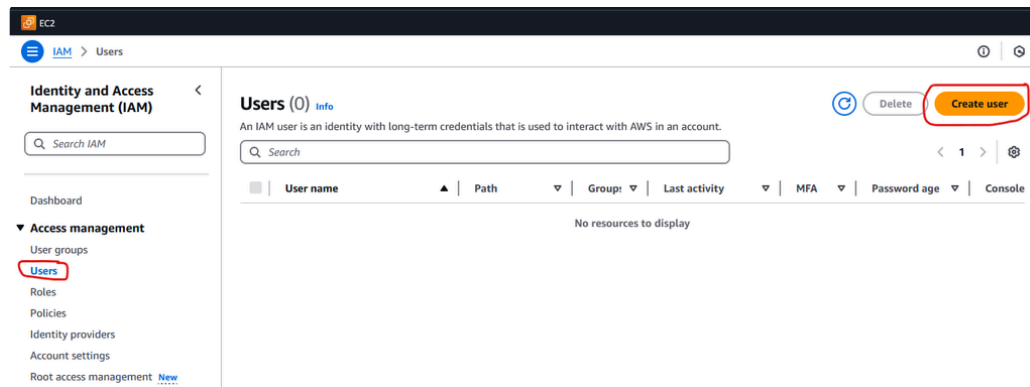
Parameter ranges

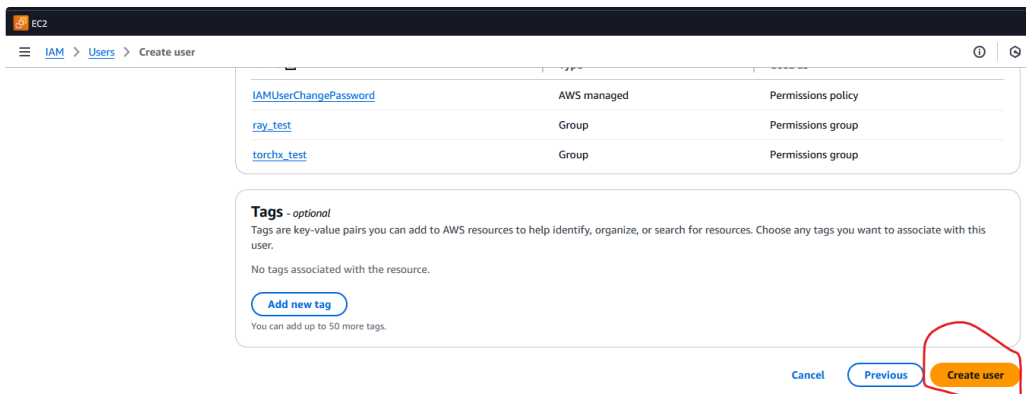
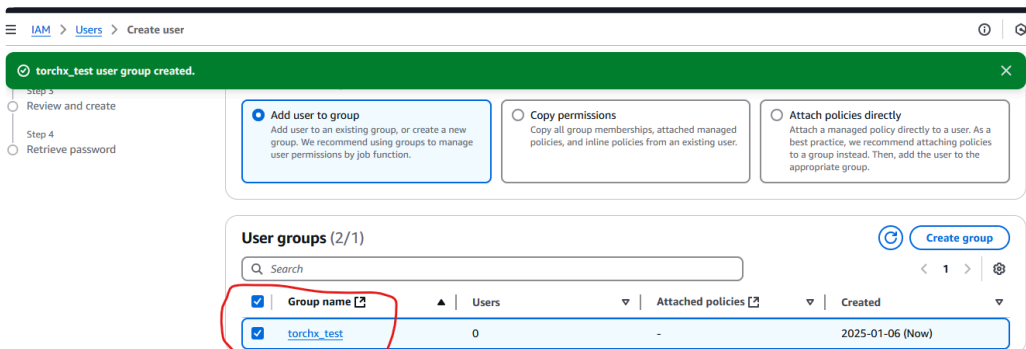
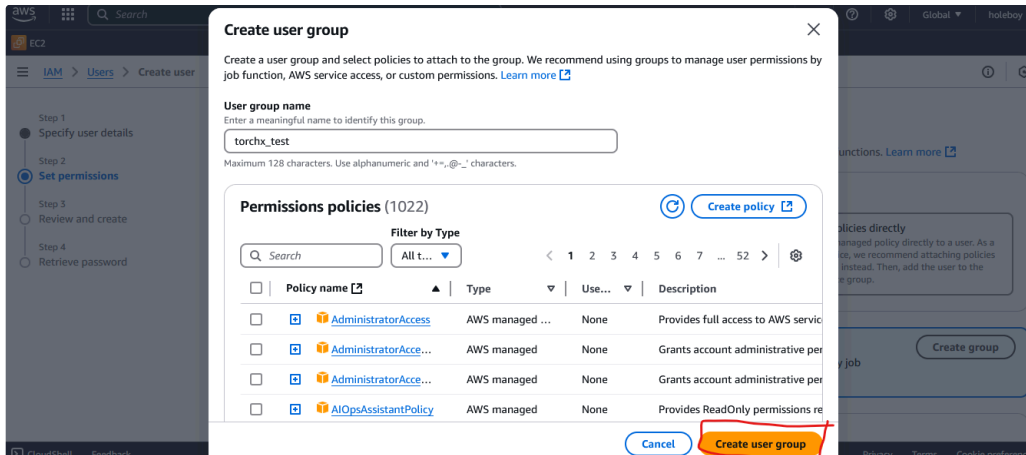
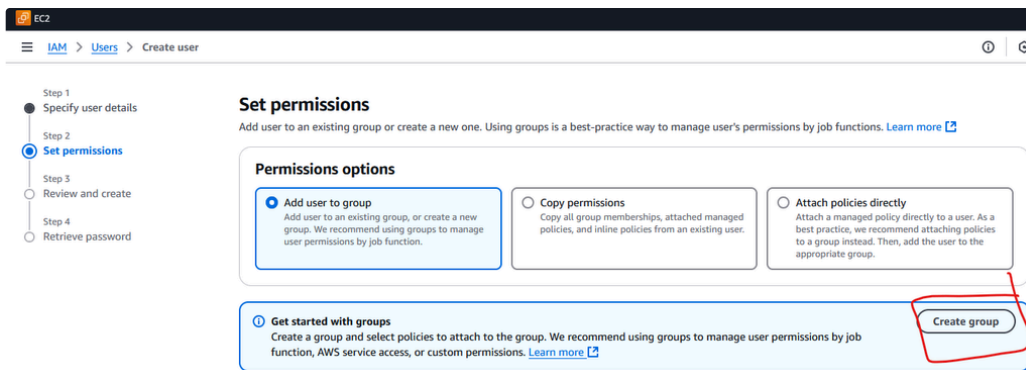
Run trials

Results

EC2 setup (free tier)

Credentials





EC2

IAM > Users

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users**
- Roles
- Policies
- Identity providers
- Account settings
- Root access management

Users (1)

An IAM user is an identity with long-term credentials that is used to interact with AWS in an account.

Search

<input type="checkbox"/>	User name	Path	Groups	Last activity	MFA	Password age	Console
<input type="checkbox"/>	<u>torchx_test</u>	/	1	-	-	0	-

EC2

IAM > Users > torchx_test

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users**
- Roles
- Policies
- Identity providers
- Account settings
- Root access management

torchx_test

Summary

ARN: [arn:aws:iam::794038232598:user/torchx_test](#)

Console access: [Enabled without MFA](#)

Created: January 06, 2025, 23:17 (UTC+13:00)

Last console sign-in: [Never](#)

Access key 1: [Create access key](#)

Permissions policies (1)

Permissions are defined by policies attached to the user directly or through groups.

Remove Add permissions

EC2

IAM > Users > torchx_test

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users**
- Roles
- Policies
- Identity providers
- Account settings
- Root access management

torchx_test

Summary

ARN: [arn:aws:iam::794038232598:user/torchx_test](#)

Console access: [Enabled without MFA](#)

Created: January 06, 2025, 23:17 (UTC+13:00)

Last console sign-in: [Never](#)

Access key 1: [Never used. Created today.](#)

Access key 2: [Create access key](#)

Permissions Groups (1) Tags (1) **Security credentials** Last Accessed

Console sign-in

Console sign-in link Console password

Manage console access

Create job queue

EC2

AWS Batch > Job queues

AWS Batch

Dashboard

Jobs

Job definitions

Job queues

Compute environments

Scheduling policies

Wizard

Console settings

- Permissions
- Jobs
- Job definition
- Job queues
- Compute environment
- Scheduling policies

Job queues (1)

Find job queue

<input type="radio"/>	Name	Type	State	Status	Priority
<input type="radio"/>	<u>test_queue</u>	ECS	Enabled	Valid	1

EC2

AWS Batch > Job queues > Create job queue

Create job queue [Info](#)

Orchestration type

☐ Fargate

☒ Amazon Elastic Compute Cloud (Amazon EC2)

☐ Amazon Elastic Kubernetes Service (Amazon EKS)

You should run your jobs on EC2 if you need access to particular instance configurations (particular processors, GPUs, or architecture) or for very-large scale workloads. Fargate jobs will start faster in the case of initial scale-out of work, as there is no need to wait for EC2 instance to launch. However, for larger workloads EC2 instances may be faster as Batch reuses instances and container images to run subsequent jobs.

We recommend that you use Amazon EC2 if your jobs require any of the following:

- More than 16 vCPUs

EC2

AWS Batch > Job queues > Create job queue

Job queue configuration

Name

torchx_queue

Job queue name must be 1-128 characters. Valid characters are a-z, A-Z, 0-9, hyphens (-), and underscores (_).

Priority [Info](#)

Job queues with a higher integer value for priority are given preference for compute environments.

1

A whole number between 0 and 1000.

Scheduling policy Amazon Resource Name (ARN) - optional [Info](#)

ARN of scheduling policy to apply to the queue. [Create scheduling policy](#)

Choose an option

Connected compute environments

When connecting compute environments to a queue, each environment must be of the same orchestration type. [Create compute environment](#)

Click to choose

A job queue must have a minimum of one and a maximum of three compute environments.

EC2

AWS Batch > Compute environments > Create compute environment

Compute environment configuration [Info](#)

Compute environment configuration

☐ Fargate

☒ Amazon Elastic Compute Cloud (Amazon EC2)

☐ Amazon Elastic Kubernetes Service (Amazon EKS)

You should run your jobs on EC2 if you need access to particular instance configurations (particular processors, GPUs, or architecture) or for very-large scale workloads. Fargate jobs will start faster in the case of initial scale-out of work, as there is no need to wait for EC2 instance to launch. However, for larger workloads EC2 instances may be faster as Batch reuses instances and container images to run subsequent jobs.

We recommend that you use Amazon EC2 if your jobs require any of the following:

EC2

AWS Batch > Compute environments > Create compute environment

Network configuration

Subnets

subnet-04ffa187e1bd181f

subnet-0acdb534772504f3d

subnet-0a20aeba98265b9ff

Security groups

sg-0b802b8d3c6faabe8

Placement group

-

► **Create compute environment request JSON** [Copy](#)

[Cancel](#) [Previous](#) [Create compute environment](#)

AWS Batch

Dashboard
Jobs
Job definitions
Job queues
Compute environments
Scheduling policies
Wizard

Console settings
Permissions
Jobs
Job definition
Job queues
Compute environment
Scheduling policies

Priority [Info](#)
Job queues with a higher integer value for priority are given preference for compute environments.
1
A whole number between 0 and 1000.

Scheduling policy Amazon Resource Name (ARN) - optional [Info](#)
ARN of scheduling policy to apply to the queue. [Create scheduling policy](#)
Choose an option

Connected compute environments
When connecting compute environments to a queue, each environment must be of the same orchestration type. [Create compute environment](#)
Click to choose
torchx_test EC2 X
Type: MANAGED State: ENABLED Status: VALID
A job queue must have a minimum of one and a maximum of three compute environments.

Compute environment order

Set user permissions

EC2

IAM > User groups

Identity and Access Management (IAM)

Search IAM

Dashboard
Access management
User groups
Users
Roles
Policies
Identity providers
Account settings
Root access management [New](#)

User groups (1) [Info](#)
A user group is a collection of IAM users. Use groups to specify permissions for a collection of users.

Search

<input type="checkbox"/>	Group name	Users	Permissions	Creation time
<input type="checkbox"/>	torchx_test	1	Not defined	43 minutes ago

EC2

IAM > User groups > torchx_test

Identity and Access Management (IAM)

Search IAM

Dashboard
Access management
User groups
Users
Roles
Policies
Identity providers
Account settings
Root access management [New](#)

Users (1) **Permissions** **Access Advisor**

Permissions policies (0) [Info](#)
You can attach up to 10 managed policies.

Simulate Remove **Add permissions**

Search

Filter by Type
All types

<input type="checkbox"/>	Policy name	Type	Attached entities
No resources to display			

EC2

IAM > User groups > torchx_test > Add permissions

Other permission policies (1/1020)
You can attach up to 10 managed policies to this user group. All of the users in this group inherit the attached permissions.

Search: batch Filter by Type: All types 4 matches

<input type="checkbox"/>	Policy name	Type	Used as	Description
<input type="checkbox"/>	AmazonMachineLearningBatch...	AWS managed	None	Grants users permission to request Am...
<input checked="" type="checkbox"/>	AWSBatchFullAccess	AWS managed	None	Provides full access for AWS Batch res...
<input type="checkbox"/>	AWSBatchServiceEventTargetR...	AWS managed	None	Policy to enable CloudWatch Event Tar...
<input type="checkbox"/>	AWSBatchServiceRole	AWS managed	Permissions policy (1)	Policy for AWS Batch service role whic...

Cancel **Attach policies**

Add AWSBatchFullAccess

Python code

Installation

```
1 poetry add torchx[kubernetes] botorch gpytorch ax-platform tensorboard pytorch-lightning torchvision boto3
  botocore torch
```

mnist.py

▼ mnist.py

```
1
2 """
3 Example training code for ``ax_multiobjective_nas_tutorial.py``
4 """
5
6 import argparse
7 import logging
8 import os
9 import sys
10 import time
11 import warnings
12
13 import torch
14 from IPython.utils import io
15 from pytorch_lightning import LightningModule, Trainer
16 from pytorch_lightning import loggers as pl_loggers
17 from torch import nn
18 from torch.nn import functional as F
19 from torch.utils.data import DataLoader
20 from torchmetrics.functional.classification.accuracy import multiclass_accuracy
21 from torchvision import transforms
22 from torchvision.datasets import MNIST
23
24 warnings.filterwarnings("ignore") # Disable data logger warnings
25 logging.getLogger("pytorch_lightning").setLevel(logging.ERROR) # Disable GPU/TPU prints
26
27 def parse_args():
28     parser = argparse.ArgumentParser(description="train mnist")
29     parser.add_argument(
30         "--log_path", type=str, required=False,
31         help="dir to place tensorboard logs from all trials",
32         default="/tmp/mnist"
33     )
34     parser.add_argument(
35         "--hidden_size_1", type=int, required=False,
36         help="hidden size layer 1", default=16
37     )
38     parser.add_argument(
39         "--hidden_size_2", type=int, required=False,
40         help="hidden size layer 2", default=16
41     )
42     parser.add_argument(
43         "--learning_rate", type=float, required=False,
44         help="learning rate", default=1e-2
45     )
46     parser.add_argument(
```

```

47     "--epochs", type=int, required=False,
48     help="number of epochs", default=1
49 )
50 parser.add_argument(
51     "--dropout", type=float, required=False,
52     help="dropout probability", default=0.0
53 )
54 parser.add_argument(
55     "--batch_size", type=int, required=False,
56     help="batch size", default=32
57 )
58 return parser.parse_args()
59
60 args = parse_args()
61
62 PATH_DATASETS = os.environ.get("PATH_DATASETS", ".")
63
64
65 class MnistModel(LightningModule):
66     def __init__(self):
67         super().__init__()
68
69         # Tunable parameters
70         self.hidden_size_1 = args.hidden_size_1
71         self.hidden_size_2 = args.hidden_size_2
72         self.learning_rate = args.learning_rate
73         self.dropout = args.dropout
74         self.batch_size = args.batch_size
75
76         # Set class attributes
77         self.data_dir = PATH_DATASETS
78
79         # Hardcode some dataset specific attributes
80         self.num_classes = 10
81         self.dims = (1, 28, 28)
82         channels, width, height = self.dims
83         self.transform = transforms.Compose(
84             [
85                 transforms.ToTensor(),
86                 transforms.Normalize((0.1307,), (0.3081,)),
87             ]
88         )
89
90         # Create a PyTorch model
91         layers = [nn.Flatten()]
92         width = channels * width * height
93         hidden_layers = [self.hidden_size_1, self.hidden_size_2]
94         num_params = 0
95         for hidden_size in hidden_layers:
96             if hidden_size > 0:
97                 layers.append(nn.Linear(width, hidden_size))
98                 layers.append(nn.ReLU())
99                 layers.append(nn.Dropout(self.dropout))
100                 num_params += width * hidden_size
101                 width = hidden_size
102         layers.append(nn.Linear(width, self.num_classes))
103         num_params += width * self.num_classes
104

```

```

105         # Save the model and parameter counts
106         self.num_params = num_params
107         self.model = nn.Sequential(*layers) # No need to use Relu for the last layer
108
109     def forward(self, x):
110         x = self.model(x)
111         return F.log_softmax(x, dim=1)
112
113     def training_step(self, batch, batch_idx):
114         x, y = batch
115         logits = self(x)
116         loss = F.nll_loss(logits, y)
117         return loss
118
119     def validation_step(self, batch, batch_idx):
120         x, y = batch
121         logits = self(x)
122         loss = F.nll_loss(logits, y)
123         preds = torch.argmax(logits, dim=1)
124         acc = multiclass_accuracy(preds, y, num_classes=self.num_classes)
125         self.log("val_acc", acc, prog_bar=False)
126         return loss
127
128     def configure_optimizers(self):
129         optimizer = torch.optim.Adam(self.parameters(), lr=self.learning_rate)
130         return optimizer
131
132     def prepare_data(self):
133         MNIST(self.data_dir, train=True, download=True)
134         MNIST(self.data_dir, train=False, download=True)
135
136     def setup(self, stage=None):
137         self.mnist_train = MNIST(self.data_dir, train=True, transform=self.transform)
138         self.mnist_val = MNIST(self.data_dir, train=False, transform=self.transform)
139
140     def train_dataloader(self):
141         return DataLoader(self.mnist_train, batch_size=self.batch_size)
142
143     def val_dataloader(self):
144         return DataLoader(self.mnist_val, batch_size=self.batch_size)
145
146
147 def run_training_job():
148
149     mnist_model = MnistModel()
150
151     # Initialize a trainer (don't log anything since things get so slow...)
152     trainer = Trainer(
153         logger=False,
154         max_epochs=args.epochs,
155         enable_progress_bar=False,
156         deterministic=True, # Do we want a bit of noise?
157         default_root_dir=args.log_path,
158     )
159
160     logger = pl_loggers.TensorBoardLogger(args.log_path)
161
162     print(f"Logging to path: {args.log_path}.")

```



```

163
164     # Train the model and log time <
165     start = time.time()
166     trainer.fit(model=mnist_model)
167     end = time.time()
168     train_time = end - start
169     logger.log_metrics({"train_time": end - start})
170
171     # Compute the validation accuracy once and log the score
172     with io.capture_output() as captured:
173         val_accuracy = trainer.validate()[0]["val_acc"]
174     logger.log_metrics({"val_acc": val_accuracy})
175
176     # Log the number of model parameters
177     num_params = trainer.model.num_params
178     logger.log_metrics({"num_params": num_params})
179
180     logger.save()
181
182     # Print outputs
183     print(f"train time: {train_time}, val acc: {val_accuracy}, num_params: {num_params}")
184
185
186 if __name__ == "__main__":
187     run_training_job()

```

Jupyter notebook

AWS credentials

▼ AWS credentials

```

1 import os
2
3 os.environ["AWS_ACCESS_KEY_ID"] = "foo"
4 os.environ["AWS_SECRET_ACCESS_KEY"] = "bar"
5 os.environ["AWS_DEFAULT_REGION"] = "ap-southeast-2"

```

TorchXRunner

▼ TorchXRunner

```

1 from pathlib import Path
2
3 import torchx
4
5 from torchx import specs
6 from torchx.components import utils
7
8
9 def trainer(
10     log_path: str,
11     hidden_size_1: int,
12     hidden_size_2: int,
13     learning_rate: float,
14     epochs: int,

```

```

15     dropout: float,
16     batch_size: int,
17     trial_idx: int = -1,
18 ) -> specs.AppDef:
19
20     # define the log path so we can pass it to the TorchX ``AppDef``
21     if trial_idx >= 0:
22         log_path = Path(log_path).joinpath(str(trial_idx)).absolute().as_posix()
23
24     return utils.python(
25         # command line arguments to the training script
26         "--log_path",
27         log_path,
28         "--hidden_size_1",
29         str(hidden_size_1),
30         "--hidden_size_2",
31         str(hidden_size_2),
32         "--learning_rate",
33         str(learning_rate),
34         "--epochs",
35         str(epochs),
36         "--dropout",
37         str(dropout),
38         "--batch_size",
39         str(batch_size),
40         # other config options
41         name="trainer",
42         script="/opt/mnist.py",
43         image="ghcr.io/jbris/torchx-aws-test:1.0.0",
44     )
45
46 import tempfile
47 from ax.runners.torchx import TorchXRunner
48
49 # Make a temporary dir to log our results into
50 log_dir = tempfile.mkdtemp()
51
52 scheduler = "aws_batch"
53 # scheduler="local_cwd"
54
55 ax_runner = TorchXRunner(
56     tracker_base="/tmp/",
57     component=trainer,
58     # NOTE: To launch this job on a cluster instead of locally you can
59     # specify a different scheduler and adjust arguments appropriately.
60     scheduler=scheduler,
61     component_const_params={"log_path": log_dir},
62     cfg={"queue": "torchx_queue"},
63 )

```

Parameter ranges

▼ Parameter ranges

```

1 from ax.core import (
2     ChoiceParameter,
3     ParameterType,
4     RangeParameter,

```

```

5     SearchSpace,
6 )
7
8 parameters = [
9     # NOTE: In a real-world setting, hidden_size_1 and hidden_size_2
10    # should probably be powers of 2, but in our simple example this
11    # would mean that ``num_params`` can't take on that many values, which
12    # in turn makes the Pareto frontier look pretty weird.
13    RangeParameter(
14        name="hidden_size_1",
15        lower=16,
16        upper=128,
17        parameter_type=ParameterType.INT,
18        log_scale=True,
19    ),
20    RangeParameter(
21        name="hidden_size_2",
22        lower=16,
23        upper=128,
24        parameter_type=ParameterType.INT,
25        log_scale=True,
26    ),
27    RangeParameter(
28        name="learning_rate",
29        lower=1e-4,
30        upper=1e-2,
31        parameter_type=ParameterType.FLOAT,
32        log_scale=True,
33    ),
34    RangeParameter(
35        name="epochs",
36        lower=1,
37        upper=4,
38        parameter_type=ParameterType.INT,
39    ),
40    RangeParameter(
41        name="dropout",
42        lower=0.0,
43        upper=0.5,
44        parameter_type=ParameterType.FLOAT,
45    ),
46    ChoiceParameter( # NOTE: ``ChoiceParameters`` don't require log-scale
47        name="batch_size",
48        values=[32, 64, 128, 256],
49        parameter_type=ParameterType.INT,
50        is_ordered=True,
51        sort_values=True,
52    ),
53 ]
54
55 search_space = SearchSpace(
56     parameters=parameters,
57     # NOTE: In practice, it may make sense to add a constraint
58     # hidden_size_2 <= hidden_size_1
59     parameter_constraints=[],
60 )
61
62 from ax.metrics.tensorboard import TensorboardMetric

```

```

63 from tensorboard.backend.event_processing import plugin_event_multiplexer as event_multiplexer
64
65 class MyTensorboardMetric(TensorboardMetric):
66
67     # NOTE: We need to tell the new TensorBoard metric how to get the id /
68     # file handle for the TensorBoard logs from a trial. In this case
69     # our convention is to just save a separate file per trial in
70     # the prespecified log dir.
71     def _get_event_multiplexer_for_trial(self, trial):
72         mul = event_multiplexer.EventMultiplexer(max_reload_threads=20)
73         mul.AddRunsFromDirectory(Path(log_dir).joinpath(str(trial.index)).as_posix(), None)
74         mul.Reload()
75
76         return mul
77
78     # This indicates whether the metric is queryable while the trial is
79     # still running. We don't use this in the current tutorial, but Ax
80     # utilizes this to implement trial-level early-stopping functionality.
81     @classmethod
82     def is_available_while_running(cls):
83         return False
84
85 val_acc = MyTensorboardMetric(
86     name="val_acc",
87     tag="val_acc",
88     lower_is_better=False,
89 )
90 model_num_params = MyTensorboardMetric(
91     name="num_params",
92     tag="num_params",
93     lower_is_better=True,
94 )

```

Run trials

▼ Run trials

```

1 from ax.core import MultiObjective, Objective, ObjectiveThreshold
2 from ax.core.optimization_config import MultiObjectiveOptimizationConfig
3
4
5 opt_config = MultiObjectiveOptimizationConfig(
6     objective=MultiObjective(
7         objectives=[
8             Objective(metric=val_acc, minimize=False),
9             Objective(metric=model_num_params, minimize=True),
10        ],
11    ),
12    objective_thresholds=[
13        ObjectiveThreshold(metric=val_acc, bound=0.94, relative=False),
14        ObjectiveThreshold(metric=model_num_params, bound=80_000, relative=False),
15    ],
16 )
17
18 from ax.core import Experiment
19
20 experiment = Experiment(
21     name="torchx_mnist",

```

```

22     search_space=search_space,
23     optimization_config=opt_config,
24     runner=ax_runner,
25 )
26
27 total_trials = 48 # total evaluation budget
28
29 from ax.modelbridge.dispatch_utils import choose_generation_strategy
30
31 gs = choose_generation_strategy(
32     search_space=experiment.search_space,
33     optimization_config=experiment.optimization_config,
34     num_trials=total_trials,
35 )
36
37 from ax.service.scheduler import Scheduler, SchedulerOptions
38
39 scheduler = Scheduler(
40     experiment=experiment,
41     generation_strategy=gs,
42     options=SchedulerOptions(
43         total_trials=total_trials, max_pending_trials=4
44     ),
45 )
46
47 scheduler.run_n_trials(3)

```

Results

▼ Results

```

1
2 from ax.service.utils.report_utils import exp_to_df
3
4 df = exp_to_df(experiment)
5 df.head(10)
6
7 from ax.service.utils.report_utils import _pareto_frontier_scatter_2d_plotly
8
9 _pareto_frontier_scatter_2d_plotly(experiment)

```