

AWS training: BoTorch TorchX - Neural architecture search

EC2 setup (free tier)

Credentials

Create job queue

Set user permissions

Python code

Installation

mnist.py

Jupyter notebook

AWS credentials

TorchXRunner

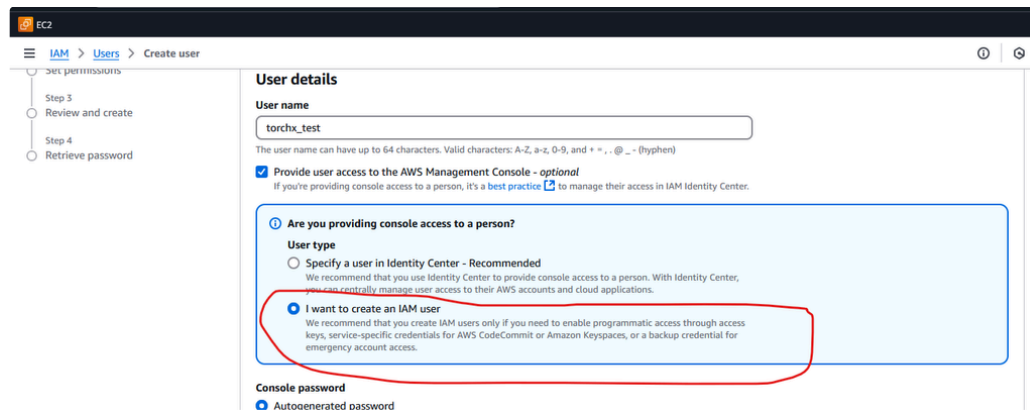
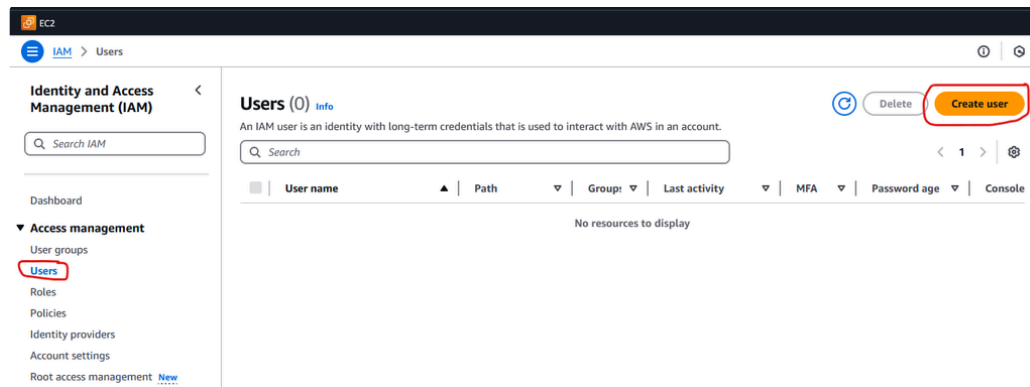
Parameter ranges

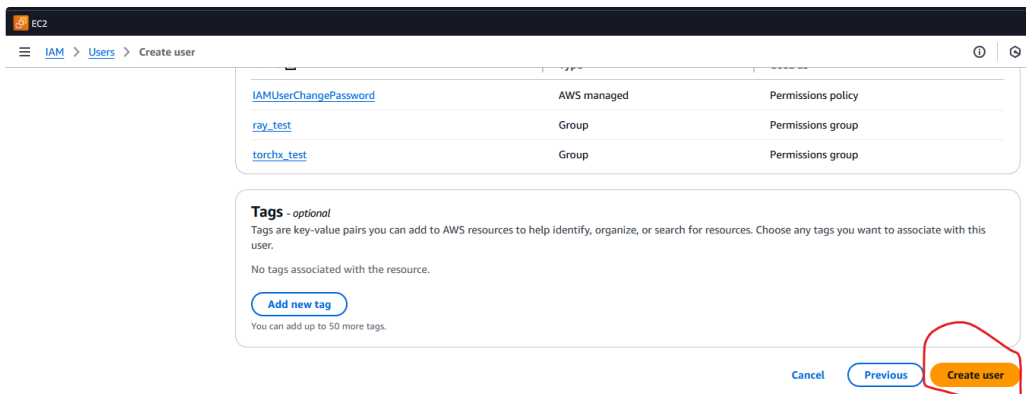
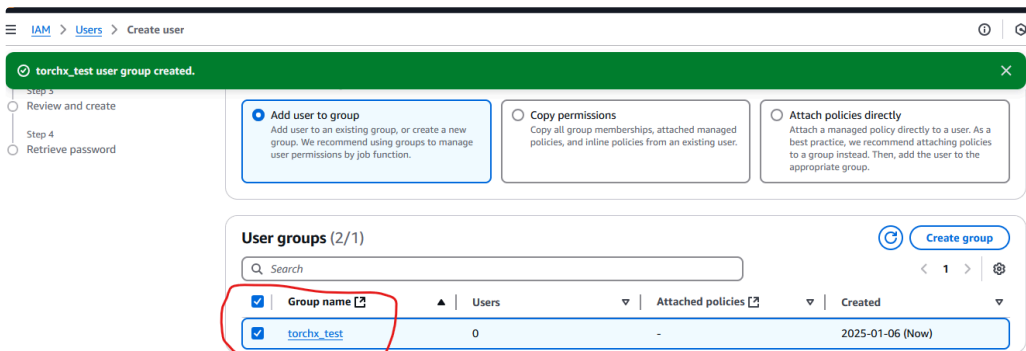
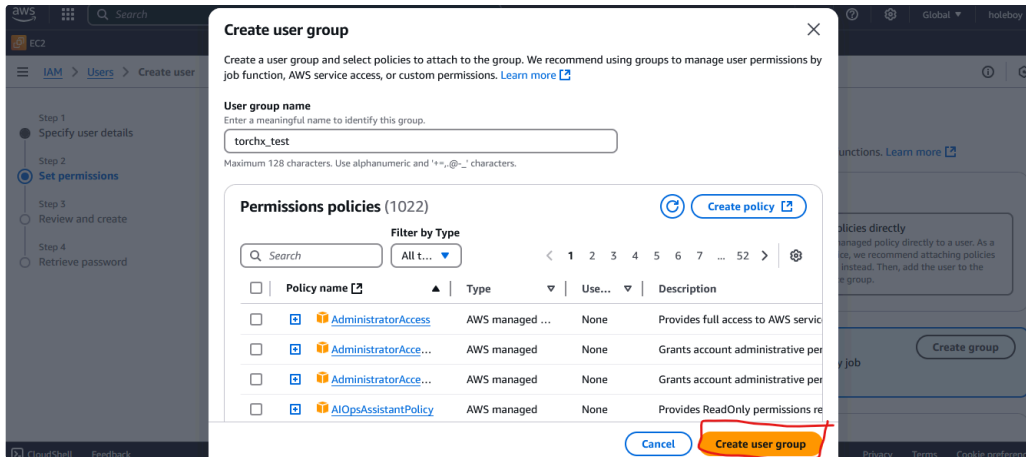
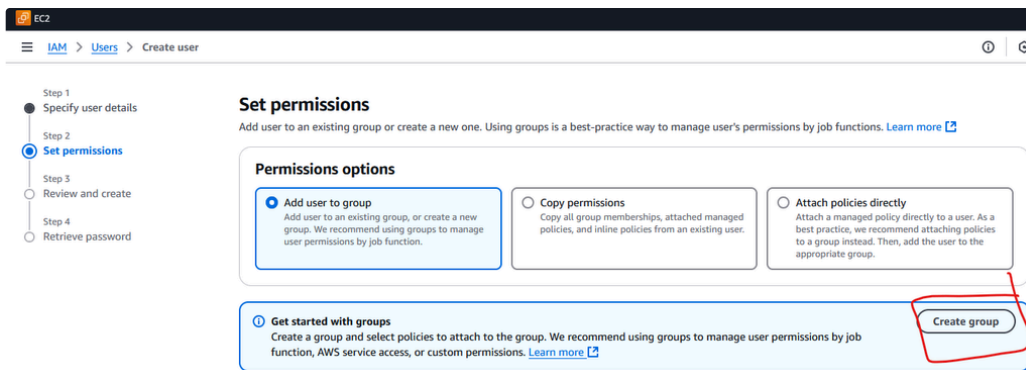
Run trials

Results

EC2 setup (free tier)

Credentials





EC2

IAM > Users

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users**
- Roles
- Policies
- Identity providers
- Account settings
- Root access management

Users (1)

An IAM user is an identity with long-term credentials that is used to interact with AWS in an account.

Search

<input type="checkbox"/>	User name	Path	Groups	Last activity	MFA	Password age	Console
<input type="checkbox"/>	<u>torchx_test</u>	/	1	-	-	0	-

EC2

IAM > Users > torchx_test

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users**
- Roles
- Policies
- Identity providers
- Account settings
- Root access management

torchx_test

Summary

ARN: [arn:aws:iam::794038232598:user/torchx_test](#)

Console access: [Enabled without MFA](#)

Created: January 06, 2025, 23:17 (UTC+13:00)

Last console sign-in: [Never](#)

Access key 1: [Create access key](#)

Permissions policies (1)

Permissions are defined by policies attached to the user directly or through groups.

Remove Add permissions

EC2

IAM > Users > torchx_test

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users**
- Roles
- Policies
- Identity providers
- Account settings
- Root access management

torchx_test

Summary

ARN: [arn:aws:iam::794038232598:user/torchx_test](#)

Console access: [Enabled without MFA](#)

Created: January 06, 2025, 23:17 (UTC+13:00)

Last console sign-in: [Never](#)

Access key 1: [Never used. Created today.](#)

Access key 2: [Create access key](#)

Permissions Groups (1) Tags (1) **Security credentials** Last Accessed

Console sign-in

Console sign-in link Console password

Manage console access

Create job queue

EC2

AWS Batch > Job queues

AWS Batch

Dashboard

Jobs

Job definitions

Job queues

Compute environments

Scheduling policies

Wizard

Console settings

- Permissions
- Jobs
- Job definition
- Job queues
- Compute environment
- Scheduling policies

Job queues (1)

Find job queue

<input type="radio"/>	Name	Type	State	Status	Priority
<input type="radio"/>	<u>test_queue</u>	ECS	Enabled	Valid	1

EC2

AWS Batch > Job queues > Create job queue

Create job queue [Info](#)

Orchestration type

☐ Fargate

☒ Amazon Elastic Compute Cloud (Amazon EC2)

☐ Amazon Elastic Kubernetes Service (Amazon EKS)

You should run your jobs on EC2 if you need access to particular instance configurations (particular processors, GPUs, or architecture) or for very-large scale workloads. Fargate jobs will start faster in the case of initial scale-out of work, as there is no need to wait for EC2 instance to launch. However, for larger workloads EC2 instances may be faster as Batch reuses instances and container images to run subsequent jobs.

We recommend that you use Amazon EC2 if your jobs require any of the following:

- More than 16 vCPUs

EC2

AWS Batch > Job queues > Create job queue

Job queue configuration

Name

torchx_queue

Job queue name must be 1-128 characters. Valid characters are a-z, A-Z, 0-9, hyphens (-), and underscores (_).

Priority [Info](#)

Job queues with a higher integer value for priority are given preference for compute environments.

1

A whole number between 0 and 1000.

Scheduling policy Amazon Resource Name (ARN) - optional [Info](#)

ARN of scheduling policy to apply to the queue. [Create scheduling policy](#)

Choose an option

Connected compute environments

When connecting compute environments to a queue, each environment must be of the same orchestration type. [Create compute environment](#)

Click to choose

A job queue must have a minimum of one and a maximum of three compute environments.

EC2

AWS Batch > Compute environments > Create compute environment

Compute environment configuration [Info](#)

Compute environment configuration

☐ Fargate

☒ Amazon Elastic Compute Cloud (Amazon EC2)

☐ Amazon Elastic Kubernetes Service (Amazon EKS)

You should run your jobs on EC2 if you need access to particular instance configurations (particular processors, GPUs, or architecture) or for very-large scale workloads. Fargate jobs will start faster in the case of initial scale-out of work, as there is no need to wait for EC2 instance to launch. However, for larger workloads EC2 instances may be faster as Batch reuses instances and container images to run subsequent jobs.

We recommend that you use Amazon EC2 if your jobs require any of the following:

EC2

AWS Batch > Compute environments > Create compute environment

Network configuration

Subnets

subnet-04ffa187e1bd181f

subnet-0acdb534772504f3d

subnet-0a20aeba98265b9ff

Security groups

sg-0b802b8d3c6faabe8

Placement group

-

► **Create compute environment request JSON** [Copy](#)

[Cancel](#) [Previous](#) [Create compute environment](#)

AWS Batch

Dashboard
Jobs
Job definitions
Job queues
Compute environments
Scheduling policies
Wizard

Console settings
Permissions
Jobs
Job definition
Job queues
Compute environment
Scheduling policies

Priority [Info](#)
Job queues with a higher integer value for priority are given preference for compute environments.
1
A whole number between 0 and 1000.

Scheduling policy Amazon Resource Name (ARN) - optional [Info](#)
ARN of scheduling policy to apply to the queue. [Create scheduling policy](#)
Choose an option

Connected compute environments
When connecting compute environments to a queue, each environment must be of the same orchestration type. [Create compute environment](#)
Click to choose
torchx_test EC2 X
Type: MANAGED State: ENABLED Status: VALID
A job queue must have a minimum of one and a maximum of three compute environments.

Compute environment order

Set user permissions

EC2

IAM > User groups

Identity and Access Management (IAM)

Dashboard
Access management
User groups
Users
Roles
Policies
Identity providers
Account settings
Root access management [New](#)

User groups (1) [Info](#)
A user group is a collection of IAM users. Use groups to specify permissions for a collection of users.

Search

<input type="checkbox"/>	Group name	Users	Permissions	Creation time
<input type="checkbox"/>	torchx_test	1	Not defined	43 minutes ago

EC2

IAM > User groups > torchx_test

Identity and Access Management (IAM)

Dashboard
Access management
User groups
Users
Roles
Policies
Identity providers
Account settings
Root access management [New](#)

Users (1) **Permissions** **Access Advisor**

Permissions policies (0) [Info](#)
You can attach up to 10 managed policies.

Search

Filter by Type: All types

<input type="checkbox"/>	Policy name	Type	Attached entities
No resources to display			

[Add permissions](#)

EC2

IAM > User groups > torchx_test > Add permissions

Other permission policies (1/1020)
You can attach up to 10 managed policies to this user group. All of the users in this group inherit the attached permissions.

Search: batch Filter by Type: All types 4 matches

<input type="checkbox"/>	Policy name	Type	Used as	Description
<input type="checkbox"/>	AmazonMachineLearningBatch...	AWS managed	None	Grants users permission to request Am...
<input checked="" type="checkbox"/>	AWSBatchFullAccess	AWS managed	None	Provides full access for AWS Batch res...
<input type="checkbox"/>	AWSBatchServiceEventTargetR...	AWS managed	None	Policy to enable CloudWatch Event Tar...
<input type="checkbox"/>	AWSBatchServiceRole	AWS managed	Permissions policy (1)	Policy for AWS Batch service role whic...

[Cancel](#) [Attach policies](#)

Add AWSBatchFullAccess

Python code

Installation

```
1 poetry add torchx[kubernetes] botorch gpytorch ax-platform tensorboard pytorch-lightning torchvision boto3
  botocore torch
```

mnist.py

▼ mnist.py

```
1
2 """
3 Example training code for ``ax_multiobjective_nas_tutorial.py``
4 """
5
6 import argparse
7 import logging
8 import os
9 import sys
10 import time
11 import warnings
12
13 import torch
14 from IPython.utils import io
15 from pytorch_lightning import LightningModule, Trainer
16 from pytorch_lightning import loggers as pl_loggers
17 from torch import nn
18 from torch.nn import functional as F
19 from torch.utils.data import DataLoader
20 from torchmetrics.functional.classification.accuracy import multiclass_accuracy
21 from torchvision import transforms
22 from torchvision.datasets import MNIST
23
24 warnings.filterwarnings("ignore") # Disable data logger warnings
25 logging.getLogger("pytorch_lightning").setLevel(logging.ERROR) # Disable GPU/TPU prints
26
27
28 def parse_args():
29     parser = argparse.ArgumentParser(description="train mnist")
30     parser.add_argument(
31         "--log_path", type=str, required=False,
32         help="dir to place tensorboard logs from all trials",
33         default="/tmp/mnist"
34     )
35     parser.add_argument(
36         "--hidden_size_1", type=int, required=False,
37         help="hidden size layer 1", default=16
38     )
39     parser.add_argument(
40         "--hidden_size_2", type=int, required=False,
41         help="hidden size layer 2", default=16
42     )
43     parser.add_argument(
44         "--learning_rate", type=float, required=False,
45         help="learning rate", default=1e-2
46     )
```

```

47     parser.add_argument(
48         "--epochs", type=int, required=False,
49         help="number of epochs", default=1
50     )
51     parser.add_argument(
52         "--dropout", type=float, required=False,
53         help="dropout probability", default=0.0
54     )
55     parser.add_argument(
56         "--batch_size", type=int, required=False,
57         help="batch size", default=32
58     )
59     return parser.parse_args()
60
61 args = parse_args()
62
63 PATH_DATASETS = os.environ.get("PATH_DATASETS", ".")
64
65
66 class MnistModel(LightningModule):
67     def __init__(self):
68         super().__init__()
69
70         # Tunable parameters
71         self.hidden_size_1 = args.hidden_size_1
72         self.hidden_size_2 = args.hidden_size_2
73         self.learning_rate = args.learning_rate
74         self.dropout = args.dropout
75         self.batch_size = args.batch_size
76
77         # Set class attributes
78         self.data_dir = PATH_DATASETS
79
80         # Hardcode some dataset specific attributes
81         self.num_classes = 10
82         self.dims = (1, 28, 28)
83         channels, width, height = self.dims
84         self.transform = transforms.Compose(
85             [
86                 transforms.ToTensor(),
87                 transforms.Normalize((0.1307,), (0.3081,)),
88             ]
89         )
90
91         # Create a PyTorch model
92         layers = [nn.Flatten()]
93         width = channels * width * height
94         hidden_layers = [self.hidden_size_1, self.hidden_size_2]
95         num_params = 0
96         for hidden_size in hidden_layers:
97             if hidden_size > 0:
98                 layers.append(nn.Linear(width, hidden_size))
99                 layers.append(nn.ReLU())
100                 layers.append(nn.Dropout(self.dropout))
101                 num_params += width * hidden_size
102                 width = hidden_size
103         layers.append(nn.Linear(width, self.num_classes))
104         num_params += width * self.num_classes

```

```

105
106     # Save the model and parameter counts
107     self.num_params = num_params
108     self.model = nn.Sequential(*layers) # No need to use Relu for the last layer
109
110     def forward(self, x):
111         x = self.model(x)
112         return F.log_softmax(x, dim=1)
113
114     def training_step(self, batch, batch_idx):
115         x, y = batch
116         logits = self(x)
117         loss = F.nll_loss(logits, y)
118         return loss
119
120     def validation_step(self, batch, batch_idx):
121         x, y = batch
122         logits = self(x)
123         loss = F.nll_loss(logits, y)
124         preds = torch.argmax(logits, dim=1)
125         acc = multiclass_accuracy(preds, y, num_classes=self.num_classes)
126         self.log("val_acc", acc, prog_bar=False)
127         return loss
128
129     def configure_optimizers(self):
130         optimizer = torch.optim.Adam(self.parameters(), lr=self.learning_rate)
131         return optimizer
132
133     def prepare_data(self):
134         MNIST(self.data_dir, train=True, download=True)
135         MNIST(self.data_dir, train=False, download=True)
136
137     def setup(self, stage=None):
138         self.mnist_train = MNIST(self.data_dir, train=True, transform=self.transform)
139         self.mnist_val = MNIST(self.data_dir, train=False, transform=self.transform)
140
141     def train_dataloader(self):
142         return DataLoader(self.mnist_train, batch_size=self.batch_size)
143
144     def val_dataloader(self):
145         return DataLoader(self.mnist_val, batch_size=self.batch_size)
146
147
148     def run_training_job():
149         import time
150
151         mnist_model = MnistModel()
152
153         # Initialize a trainer (don't log anything since things get so slow...)
154         trainer = Trainer(
155             logger=False,
156             max_epochs=args.epochs,
157             enable_progress_bar=False,
158             deterministic=True, # Do we want a bit of noise?
159             default_root_dir=args.log_path,
160         )
161
162         logger = pl_loggers.TensorBoardLogger(args.log_path)

```



```

163
164     print(f"Logging to path: {args.log_path}.")
165
166     # Train the model and log time ✂
167     start = time.time()
168     trainer.fit(model=mnist_model)
169     end = time.time()
170     train_time = end - start
171     logger.log_metrics({"train_time": end - start})
172
173     # Compute the validation accuracy once and log the score
174     with io.capture_output() as captured:
175         val_accuracy = trainer.validate()[0]["val_acc"]
176     logger.log_metrics({"val_acc": val_accuracy})
177
178     # Log the number of model parameters
179     num_params = trainer.model.num_params
180     logger.log_metrics({"num_params": num_params})
181
182     logger.save()
183
184     time.sleep(25)
185
186     # Print outputs
187     print(f"train time: {train_time}, val acc: {val_accuracy}, num_params: {num_params}")
188
189
190 if __name__ == "__main__":
191     run_training_job()

```

Jupyter notebook

AWS credentials

▼ AWS credentials

```

1 import os
2
3 os.environ["AWS_ACCESS_KEY_ID"] = "foo"
4 os.environ["AWS_SECRET_ACCESS_KEY"] = "bar"
5 os.environ["AWS_DEFAULT_REGION"] = "ap-southeast-2"

```

TorchXRunner

▼ TorchXRunner

```

1 from pathlib import Path
2
3 import torchx
4
5 from torchx import specs
6 from torchx.components import utils
7
8 src_log_dir = f"{os.getcwd()}/output"
9 dst_log_dir = "/output"
10

```

```

11 def trainer(
12     log_path: str,
13     hidden_size_1: int,
14     hidden_size_2: int,
15     learning_rate: float,
16     # epochs: int,
17     dropout: float,
18     batch_size: int,
19     trial_idx: int = -1,
20 ) -> specs.AppDef:
21
22     # define the log path so we can pass it to the TorchX ``AppDef``
23     if trial_idx >= 0:
24         log_path = Path(log_path).joinpath(str(trial_idx)).absolute().as_posix()
25
26     return utils.sh(
27         "python",
28         "mnist.py",
29         "--log_path",
30         log_path,
31         "--hidden_size_1",
32         str(hidden_size_1),
33         "--hidden_size_2",
34         str(hidden_size_2),
35         "--learning_rate",
36         str(learning_rate),
37         # "--epochs",
38         # str(epochs),
39         "--dropout",
40         str(dropout),
41         "--batch_size",
42         str(batch_size),
43         # other config options
44         # name="trainer",
45         # script="mnist.py",
46         image="ghcr.io/jbris/torchx-aws-test:1.0.0",
47         mounts=[
48             "type=bind",
49             f"src={src_log_dir}",
50             f"dst={dst_log_dir}",
51             "perm=rwm"
52         ]
53     )
54
55 import tempfile
56 from ax.runners.torchx import TorchXRunner
57
58 # Make a temporary dir to log our results into
59
60 scheduler = "aws_batch"
61 scheduler="local_cwd"
62 scheduler="local_docker"
63
64 ax_runner = TorchXRunner(
65     tracker_base="/tmp/",
66     component=trainer,
67     scheduler=scheduler,
68     component_const_params={"log_path": dst_log_dir},

```

```
69     cfg={"queue": "torchx_queue"},
70 )
```

Parameter ranges

▼ Parameter ranges

```
1  from ax.core import (
2      ChoiceParameter,
3      ParameterType,
4      RangeParameter,
5      SearchSpace,
6  )
7
8  parameters = [
9      # NOTE: In a real-world setting, hidden_size_1 and hidden_size_2
10     # should probably be powers of 2, but in our simple example this
11     # would mean that ``num_params`` can't take on that many values, which
12     # in turn makes the Pareto frontier look pretty weird.
13     RangeParameter(
14         name="hidden_size_1",
15         lower=4,
16         upper=8,
17         parameter_type=ParameterType.INT,
18         log_scale=True,
19     ),
20     RangeParameter(
21         name="hidden_size_2",
22         lower=4,
23         upper=8,
24         parameter_type=ParameterType.INT,
25         log_scale=True,
26     ),
27     RangeParameter(
28         name="learning_rate",
29         lower=1e-2,
30         upper=1e-1,
31         parameter_type=ParameterType.FLOAT,
32         log_scale=True,
33     ),
34     # RangeParameter(
35     #     name="epochs",
36     #     lower=1,
37     #     upper=2,
38     #     parameter_type=ParameterType.INT,
39     # ),
40     RangeParameter(
41         name="dropout",
42         lower=0.0,
43         upper=0.5,
44         parameter_type=ParameterType.FLOAT,
45     ),
46     ChoiceParameter( # NOTE: ``ChoiceParameters`` don't require log-scale
47         name="batch_size",
48         values=[16, 32],
49         parameter_type=ParameterType.INT,
50         is_ordered=True,
51         sort_values=True,
```

```

52     ),
53 ]
54
55 search_space = SearchSpace(
56     parameters=parameters,
57     # NOTE: In practice, it may make sense to add a constraint
58     # hidden_size_2 <= hidden_size_1
59     parameter_constraints=[],
60 )
61
62 from ax.metrics.tensorboard import TensorboardMetric
63 from tensorboard.backend.event_processing import plugin_event_multiplexer as event_multiplexer
64
65 class MyTensorboardMetric(TensorboardMetric):
66
67     # NOTE: We need to tell the new TensorBoard metric how to get the id /
68     # file handle for the TensorBoard logs from a trial. In this case
69     # our convention is to just save a separate file per trial in
70     # the prespecified log dir.
71     def _get_event_multiplexer_for_trial(self, trial):
72         mul = event_multiplexer.EventMultiplexer(max_reload_threads=20)
73         mul.AddRunsFromDirectory(Path(src_log_dir).joinpath(str(trial.index)).as_posix(), None)
74         mul.Reload()
75
76         return mul
77
78     # This indicates whether the metric is queryable while the trial is
79     # still running. We don't use this in the current tutorial, but Ax
80     # utilizes this to implement trial-level early-stopping functionality.
81     @classmethod
82     def is_available_while_running(cls):
83         return True
84
85 val_acc = MyTensorboardMetric(
86     name="val_acc",
87     tag="val_acc",
88     lower_is_better=False,
89 )
90 model_num_params = MyTensorboardMetric(
91     name="num_params",
92     tag="num_params",
93     lower_is_better=True,
94 )
95

```

Run trials

▼ Run trials

```

1 from ax.core import MultiObjective, Objective, ObjectiveThreshold
2 from ax.core.optimization_config import MultiObjectiveOptimizationConfig
3
4
5 opt_config = MultiObjectiveOptimizationConfig(
6     objective=MultiObjective(
7         objectives=[
8             Objective(metric=val_acc, minimize=False),
9             Objective(metric=model_num_params, minimize=True),

```

```

10     ],
11 ),
12 objective_thresholds=[
13     ObjectiveThreshold(metric=val_acc, bound=0.94, relative=False),
14     ObjectiveThreshold(metric=model_num_params, bound=80_000, relative=False),
15 ],
16 )
17
18 from ax.core import Experiment
19
20 experiment = Experiment(
21     name="torchx_mnist",
22     search_space=search_space,
23     optimization_config=opt_config,
24     runner=ax_runner,
25 )
26
27 total_trials = 48 # total evaluation budget
28
29 from ax.modelbridge.dispatch_utils import choose_generation_strategy
30
31 gs = choose_generation_strategy(
32     search_space=experiment.search_space,
33     optimization_config=experiment.optimization_config,
34     num_trials=total_trials,
35 )
36
37 from ax.service.scheduler import Scheduler, SchedulerOptions
38
39 scheduler = Scheduler(
40     experiment=experiment,
41     generation_strategy=gs,
42     options=SchedulerOptions(
43         total_trials=total_trials, max_pending_trials=4,
44         init_seconds_between_polls=5, seconds_between_polls_backoff_factor=1
45     ),
46 )
47
48 scheduler.run_n_trials(1)

```

Results

▼ Results

```

1
2 from ax.service.utils.report_utils import exp_to_df
3
4 df = exp_to_df(experiment)
5 df.head(10)
6
7 from ax.service.utils.report_utils import _pareto_frontier_scatter_2d_plotly
8
9 _pareto_frontier_scatter_2d_plotly(experiment)

```